Introduction to Data Access and Storage

```
$ echo "Data Sciences Institute"
```

Introduction:

Welcome

What is SQL?

Data Modelling

Welcome / About Us

About Us (Thomas)

- Started with SQL in first job after graduating

- Data Analyst in US healthcare until moving to Canada in 2018

- Transitioned to Data Engineer at Plan Canada

- Pursued Master of Information at UofT for R and Python focus

- Currently working on Data Governance and Ethics at Thomson Reuters

- ***Have a 2 year old Samoyed named Alto who takes up a lot ofmy free time*** 👉

About Us (Ananya)

- UofT Computer Science and Statistics graduate, with minor in Mathematics

- Co-op at Citi (Software Development) and BMO (Business Analytics)

- RBC Capital Markets BSA using SQL for Salesforce CRM data, transitioning to Supervisory and Regulatory role soon

- Former Women in Computer Science president at UofT

- Involved with Toronto Wxmen in Data Science and RWomen at RBC

About Us (Vishnou)

- Paris-Saclay University grad in Applied Math/Data Science

- Explored SQL, Python, R in academia, then applied in industry

- ML internships at Docugami, Ekimetrics, Stellantis, Randstad

- Currently in UofT MscAC AI program

- Working on Data Science/LLMs intersection projects

- Considering PhD for professorship

- Hobbies: guitar, saxophone, bass in music bands 🎸 🎷

Welcome / Course Content

Course Content

- Data Modelling, Data Structures, Schemas, Basic Data Management, Normal Forms

- Basic SQL Syntax

- Essential SQL Syntax

- Advanced Techniques

- Importing and Exporting Data to and from SQL

- SQL's relationship to the Machine Learning Pipeline

- Data Stewardship, Ethics, and SQL in the Wild

GitHub Repo

https://github.com/UofT-DSI/sql

- Schedule

- These slides (HTML & PDF)

- Our database for live coding

- All in-class code

- Homework (and answers, posted afterwards)

- Assignment details and rubrics

- Due Dates

It is crucial you visit the repo throughout the course, as I may make minor changes and push new content.

Course Content

- This course is an *Introduction* to SQL

- At the end of the course, I hope you will:

  - Feel comfortable with SQL

  - Know how to search for the right thing on Stack Overflow

  - Read documentation

- We won't cover advanced topics like:

  - Stored Procedures, Triggers, Jobs

  - DBA work (monitoring, server setup, etc)

  - Complex ETL or tooling

Homework (required)

- Six homeworks total

- At the end of each SQL module, I will provide you with 3-10 queries to write on your own
  - Covers the topics of the modules

- Review answers in Tutorials with Ananya and Vishnou
  - I may go over one or two queries at the start of class if desired

- Designed to be relatively easy
  - Reaffirms what we wrote together
  - Doing work on your own helps reinforce the learning

- ChatGPT probably won't help you much

Grading

- Pass/Fail. Do the work, pass the course :)

- One Assignment: Data Model Design, 40%, due TBD

- Homework: 48% (8% each), due on Thursdays and Saturdays

- Class Attendance: 12% (2% each)

  - Let Ananya or myself know if you are unable to attend a lesson

  - Code along!! Best way to learn.

What questions do you have about the course?

Welcome / Quick Technical Check

Quick Technical Check

Let's make sure everyone has DB Browser for SQLite installed.

- If not, please download it here: https://sqlitebrowser.org/dl/

For live coding:

- Please download/fork the FarmersMarket.db from our GH repo:
  - FarmersMarket.db
- Open it in SQLite with the "Open Database" button and navigate to wherever you have saved it

Good to go?

If not, please message Ananya or Vishnou

Inrtoduction:

Welcome

What is SQL?

Data Modelling

What is SQL? / SQL

- SQL: Structured Query Language

  - Pronounced as either S.Q.L. (ess-cue-ell) or "sequel"

- SQL is a *query* language rather than a programming language

  - Querying is closer to telling a computer *what you want*, rather than *what is has to do*

  - SQL code is often less reproducible than other programming languages because it's domain specific

    - Some SQL code, especially more advanced procedural code, is reproducible within the same flavour

  - SQL's domain is databases and is based on set theory

- Designed to manage data within Relational Database Management Systems (RDBMs), e.g.

  - MSSQL

  - Oracle DB

SQL Formatting

- Like other programming/query languages, SQL has reserved keywords/commands to perform instructional operations

  - Generally, these keywords are written in all caps: `SELECT`

  - Most modern interpreters no longer require this, but it is the expected standard

- All statements/queries should end with a semicolon

  - A few SQL constructs (like common table expressions, we'll get to these later) require them, otherwise they are optional

    - I'll almost certainly forget to use them

    - There's some debate over whether or not it's best practice

SQL Formatting

- In SQL, white space and/or line breaks do not matter

    - Readability is important

        - Try to keep SQL statements to a reasonable screen width

        - Use sensible line breaks

        - Offset subqueries with indents

- Code is commented in/out with `--` rather than `#`

- Code blocks can be commented out with `/* */`

```
/*

somecode spanning
multiple lines

*/
```

What is SQL? / Flavours

- RDBMs differ from one to the next:

  - different keywords

    - e.g. return only 10 rows:

      `SELECT TOP 10...` vs `SELECT ... LIMIT 10`

  - different syntax

    - e.g. not equal:

      `!=` or `<>` (or both)

  - other, more nuanced/complex differences

    - e.g. architecture, data types, etc

- We are using SQLite:

  - Super easy to get setup

  - Requires almost no overhead

  - Open source, *free*

  - Similar enough in syntax to learn on

- Broad observations about Open Source systems:

  - Excellent at what they are designed for

  - Varying data types (SQLite has some unique ones!)

  - Not every command exists, but workarounds are usually possible

  - Some utilize RDBMs that feel extremely outdated

- Broad observations about enterprise systems:

  - Powerful and designed to handle edge cases

  - Feel a bit more refined
    - Can be version dependent

  - Tend to "lock in" businesses/organizations
    - Migration is costly, sometimes outrageously so

  - Newer players (Snowflake, Databricks, etc) and cloud providers (Azure, AWS, GCP, etc) offer a lot more functionality than just database querying

What is SQL? / Environments for SQL

Databases

- Relational databases are a collection of tables, views, procedural code, and other SQL-assisting artefacts
  - Generally the data stored in a database will be related to a real-world concept
  - Backends to data-collecting systems are often databases
    - e.g. CRMs, EMR software, ERPs, web-based applications
  - Usually not connected to other databases unless deemed necessary
  - Often transactional, meaning data is actively being written to by frontend systems
  - Tables are normalized
- There are also non-relational databases, often referred to as NoSQL
  - We won't cover these
  - Common tools include: Amazon DynamoDB, Azure CosmosDB, MongoDB, Google Cloud Datastore

Data Warehouses and Data Marts

- Data Warehouses are highly structured collections of (usually tabular) data
  - Data has been processed for a specific purpose, e.g. analytics
  - Data has been centralized
    - Often with the assistance of ETL (Extract, Transform, Load) tools
  - Have a lot of overhead, require governance, and strict rigidity
  - Tables are denormalized
  - Very common for enterprises, but losing traction in many industries
- Data Marts are created from Data Warehouses to focus on a single subject
  - Designed to make Data Warehouses easier to use
  - Data is structured, but flexibility is driven by the purpose of the Data Mart
  - Some denormalized tables may be normalized or undergo even greater normalization
    - *The subject/purpose of the Data Mart might drive these types of decisions*
  - Common for enterprises that have Data Warehouses

Data Lakes and Data Swamps

- Data Lakes allow on-demand access to raw, semi-structured, structured, and unstructured data
  - Not defined by a specific purpose
  - Highly scalable
  - Can be transactional, if systems are designed to produce outputs into Data Lakes
  - Often data sources for machine learning pipelines live in Data Lakes
  - Inexpensive compared to Data Warehouses
  - Increasingly common for enterprises to shift towards Data Lakes, especially with support from newer tools like Snowflake, Databricks, etc, which can maximize the analytical value of a Data Lake
- Data Swamps...are poorly governed Data Lakes
  - Lack of documentation, lack of governance, poorly designed Data Lakes become Swamps
  - Avoid building these

What questions do you have?

Inrtoduction:

Welcome

What is SQL?

Data Modelling

Data Modelling Topics (its a lot!)

- Relational Database Management Systems

- Data Models

- Structure of Data

- Constraints

- Entity Relationship Diagrams

- Attributes of an ERD: Entities & Relationships

- Relationship Examples

- Conceptual, Logical, Physical Models

- Assignment 1: Design a Logical Model

Relational Database Management Systems

- Relational Database Management Systems (RDBMs) are software designed to:

  - Store large amounts of data

  - Utilize a query language to allow easy retrieval of the data

  - Allow multiple users to access the data simultaneously

  - Manage permissions for data access

  - Mitigate data corruption and unauthorized access

- Generally, data is stored in a *database*

  - A database is a collection of information

  - Within a database, a collection of objects (e.g. tabular data "tables") is stored

- RDBMs allow users to define interactions between these objects, such as:

  - Establish the relationship between objects

  - Define procedural scripts to query specific data or trigger an action

31

Data Models

- A data model is a notation for describing data or information

- Data models consist of:

    - Structure of the data

    - Operations

    - Constraints on the data

    - Relationships

Structure of Data

- SQL is comprised of tables

		Coat				

- Tables have Attributes and Observations

  - In SQL we call Attributes "Columns"
    - e.g. Breed, Coat Type, Coat Length

  - and Observations "Rows"
    - e.g. Samoyed, Double, Long

- SQL databases require tables to be named

  - e.g. We can call this table "breed_traits"

- Common ones include:
  - `INT` (integers: 1,2,3,-1,-2,-3)
    - most systems conserve storage space specifying their range
  - `FLOAT`, `DECIMAL`, `REAL` (decimal: 5.5, 3.333333)
  - `VARCHAR`, `NVARCHAR`, `TEXT` (text strings, with a maximum length associated: 'abc')
  - `DATE`, `DATETIME`, `TIME` (dates and times: '2023-01-09', '11:11:11.000')
- These may vary slightly by flavour (in SQLite they are simpler and less restricted)
- Data types are important:
  - They affect operation speed, storage size, data validity
    - Speed: it's computationally less expensive to compute smaller values
    - Storage: small is usually better, but the wrong size will affect systems; e.g. Psy's Gangnam Style exceeded 2,147,483,647 (32-bit signed, 2^32/2-1) views, causing YouTube to expand the view counter to 64-bit (~9.2b if

Constraints

- Data Models also specify constraints

- Constraints are rules that must be followed:
  - Referential-Integrity constraints
    - Ensure that values in one table have corresponding values in another table
  - Attribute Constraints
    - Ensure that certain types of values are always consistent within columns
    - May also ensure whether values are unique, not missing, etc

- If a value can be missing or not

- **UNIQUE**

  - All values are different

- **PRIMARY KEY (PK)**

  - Ensures each value in a column is unique within the table (e.g. an ID field)

  - One PK per table

  - Cannot be NULL

  - Ensures database integrity by restricting record deletion

- **FOREIGN KEY (FK)**

  - Creates a linkage between a column in one table and a column in another table
    - Generally, foreign keys are linked to primary keys
    - Sometimes share the same name as the linked column, but this isn't required
    - Linkage requires data types to be the same

Entity Relationship Diagrams

- Entity Relationship Diagrams (ERDs) are diagrams depicting the structure of tables within a database

    - This both *identifies the tables* and *describes their relationships*

- ERDs are useful for:

    - Database design

    - Debugging

    - Writing logical, consistent, and efficient queries

- There are three levels of detail for ERD depictions:

    - Conceptual model

    - Logical model

    - Physical model

Attributes of an ERD Entity

- For a given table:
  - Name
  - Relationship to another table
  - Column Names
  - Column Types
  - Primary Keys (if present)
  - Foreign Keys (if present)

Attributes of an ERD Relationship

- Defines which columns are related

- Defines what type of relationship exists:

    - One-to-One

    - One-to-Many

    - Many-to-Many

Relationship Examples

One-to-One: where a given row within a table is associated with only a single row in another table

Table 1: Country — Table 2: Capital City

Table 1:Country		Table 2:Capital
Canada	1:1	Ottawa
USA	1:1	Washington DC
Mexico	1:1	Mexico City

Relationship Examples

One-to-Many: where a given row within a table can be referenced by multiple rows in another table

Table 1: Country — Table 2: States

Table 1:Country		Table 2:States
Canada	1:∞	Alberta
Canada	1:∞	British Columbia
Canada	1:∞	...(11 more rows)
USA	1:∞	Alabama
USA	1:∞	Alaska

Relationship Examples

Many-to-Many: where multiple rows within a table can be referenced by multiple rows in another table

Table 1: Employee — Table 2: Employee Type (Next slide)

For this example, imagine the following:

Previously, A was an course instructor at DSI (python) *and* a TA at DSI (SQL). Thomas remains the instructor for SQL. Ananya is now the course support. If a table contained Employee ID and Employee Type ID, A would be listed twice.

Table 1: Employee — Table 2: Employee Type

Table 3: listed twice

We'd create additional Many-to-Many relationships if we created a course table, because A and Ananya are involved in multiple courses!

relationships

- Our Farmers Market database:
  - 10 tables
  - Relationships between these tables
    - e.g. product and product_category: *what type of thing a product is*
    - product and customer_purchases: *what products a customer has bought*
    - product and vendor_inventory: *what products each vendor has available*

- Often indicate the type of relationship
    - One-to-One
    - One-to-Many
    - Many-to-Many
- Our (partial) Farmers Market database:
    - product (5 columns) shares a One-to-Many relationship with vendor_inventory (5 columns) on product_id
    - product_category (2 columns) shares a One-to-Many relationship with

- Add additional detail to the logical model by adding key type and column data type

- Our (partial) Farmers Market database:
  - vendor_id (int) is the PK for vendor, which shares a One-to-Many relationship with vendor_inventory on vendor_id (FK)
  - product_id (int) is a FK for vendor_inventory (*so elsewhere in this diagram, we'd connect this to a PK of another table*)
  - market_date (date) is the PK for vendor_inventory

What questions do you have about Data Modelling?

Assignment 1

Q1) Create a logical model for a small bookstore. 📚

At the minimum it should have employee, order, sales, customer, and book entities (tables). Determine sensible column and table design based on what you know about these concepts. Keep it simple, but work out sensible relationships to keep tables reasonably sized. Include a date table. There are several tools online you can use, I'd recommend Lucidchart or Draw.io.

Q2) We want to create employee shifts, splitting up the day into morning and evening. Add this to the ERD.

Q3) The store wants to keep customer addresses. Propose two architectures for the CUSTOMER_ADDRESS table, one that will retain changes, and another that will overwrite. Which is type 1, which is type 2? *Hint, search type 1 vs type 2 slowly changing dimensions*.

Bonus: Are there privacy implications to this, why or why not?

Q4) Review the AdventureWorks ERD here: link

What questions do you have about anything from today?